

EXPERIENCE

Lead Applied Scientist

Apr 2024–Present

Future AGI — AI evaluation & observability platform for enterprise LLM systems

India

- **Research leadership** — Founding team member leading and mentoring 4–6 researchers and engineers working on evaluation, guardrails, and synthetic data. Set research direction, assign projects, run performance reviews, and partner with engineering and product to put the research into production. Co-inventor on **2 granted US patents** [10], [11].
- **Protect** — Led the team behind Protect, covering data curation and teacher-assisted relabeling, adapter design, multi-GPU training, ablations, evaluation, and deployment. Fine-tuned Gemma-3n-E4B with category-specific LoRA adapters on 113K text, image, and audio samples using 8×H100 GPUs. Outperformed GPT-4.1, WildGuard, and LlamaGuard-4 across all four safety categories; prompt-injection F1 was 96.6 versus 87.0 for the next best model, with a 65 ms median time-to-label [1].
- **AgentCompass** — Built the agent-evaluation system used on live customer workloads. On TRAIL error localization, it reached 0.657 versus 0.546 for Gemini-2.5-Pro, and found real failures that the benchmark’s annotators had missed [2].
- **Platform** — Built and shipped the production stack for multimodal evaluation, error localization, synthetic data, golden datasets, RAG, agentic workflows, and feedback loops. Extended it to voice with a call simulator and *Fix My Agent*, which tells developers what broke in their assistant and why.
- **Customers & communication** — Lead discovery, solution design, demos, and pilots for the largest enterprise accounts, and bring recurring customer problems into the product roadmap. Host the company webinar series (10+ sessions, averaging 300 registrations and 150+ attendees per session), write engineering blog posts, and own developer-facing design and open-source contribution policy.

Research Fellow

Jul 2022–Apr 2024

Microsoft Research

India

Supervisors: Dr. Kalika Bali, Dr. Sunayana Sitaram

- **Multilingual evaluation** — Built and calibrated an LLM-as-judge framework using 20K native-speaker judgments across 8 languages. **6 Microsoft Copilot teams** adopted it as a pre-launch quality gate, replacing weeks of manual review. Published the first end-to-end protocol for meta-evaluating LLM evaluators in 10 languages [3], [4].
- **MEGA** — Owned question answering and responsible-AI evaluation (toxicity, gender bias in translation) for the first broad multilingual benchmark of generative LLMs, covering 16 datasets and 70 languages. Now at **520+ citations** [6].
- **Shiksha Co-pilot** — Designed and ran harms measurement across Microsoft’s mandatory RAI taxonomy for a teacher-facing GenAI product. The product passed deployment-safety and sensitive-use reviews and now serves Karnataka government schools, cutting lesson planning from roughly 90 minutes to 90 seconds.
- **Sanmati** — Won **Best Paper at FAccT’24** for work on gender bias in Hindi language technology [5]. Also built the first normative-rating dataset for gender bias in GPT output, scoring severity by degree instead of using a binary label [8].

Researcher

Feb 2021–Jun 2022

Explainable & Reliable AI (ERAI) Research Group

Maastricht, Netherlands

PI: Prof. Nava Tintarev — Twitter-funded research appointment

- **Viewpoint diversity** — Defined two continuous measures of viewpoint diversity in political conversations, moving beyond binary “echo chamber” labels [7]. WSDM’23.

- **Ruddit** — Built the first Reddit dataset scored for offensiveness on a continuous -1 to 1 scale, using Best–Worst Scaling to avoid the rating-scale biases in earlier corpora [9]. The dataset and method were later taken up in Google Jigsaw's toxic-severity rating challenge. Facebook Online Safety Benchmark Research Award; ACL'21.

EDUCATION

- **MSc in Artificial Intelligence**, University of Amsterdam, Netherlands 2018–2020
- **BTech in Computer Science & Engineering**, Manipal University Jaipur, India 2014–2018

TECHNICAL SKILLS

- **Evaluation & Safety:** LLM-as-judge design and calibration, meta-evaluation, benchmark construction, guardrails, agentic-workflow evaluation, multilingual and multimodal evaluation, responsible-AI harms measurement, human annotation design (Best–Worst Scaling)
- **Model Training & ML:** PyTorch, Hugging Face Transformers, Axolotl, PEFT/LoRA, TRL, multi-GPU fine-tuning, ablation design, vLLM, Scikit-learn, Pandas
- **Production LLM Systems:** RAG (LangChain, LlamaIndex, vector databases, embeddings), agent frameworks (LangGraph, MCP), multi-agent orchestration
- **Languages:** Python, SQL, C, C++, Java

SELECTED PUBLICATIONS

1,050+ citations • h-index 12 • i10-index 13

(Google Scholar, Aug 2026)

- [1] K. Avinash, N. Pareek, and **R. Hada**, *Protect: Towards Robust Guardrail Stack for Trustworthy Enterprise LLM Systems*, 2025. arXiv: 2510.13351 [cs.CL].
- [2] N. V. J. K. Kartik, G. Sapra, **R. Hada**, and N. Pareek, *AgentCompass: Towards Reliable Evaluation of Agentic Workflows in Production*, 2025. arXiv: 2509.14647 [cs.AI].
- [3] **R. Hada**, V. Gumma, M. Ahmed, K. Bali, and S. Sitaram, “METAL: Towards Multilingual Meta-Evaluation”, in *Findings of the ACL: NAACL 2024*, 2024. DOI: 10.18653/v1/2024.findings-naacl.148
- [4] **R. Hada**, V. Gumma, A. Wynter, H. Diddee, M. Ahmed, M. Choudhury, K. Bali, and S. Sitaram, “Are Large Language Model-based Evaluators the Solution to Scaling Up Multilingual Evaluation?”, in *Findings of the ACL: EACL 2024*, 2024.
- [5] **R. Hada**, S. Husain, V. Gumma, H. Diddee, A. Yadavalli, A. Seth, N. Kulkarni, U. Gadiraju, A. Vashistha, V. Seshadri, and K. Bali, “Akal Badi ya Bias: An Exploratory Study of Gender Bias in Hindi Language Technology”, in *Proceedings of the 2024 ACM Conference FAccT '24*, 2024. DOI: 10.1145/3630106.3659017
- [6] K. Ahuja, H. Diddee, **R. Hada**, M. Ochieng, K. Ramesh, P. Jain, A. Nambi, T. Ganu, S. Segal, M. Ahmed, K. Bali, and S. Sitaram, “MEGA: Multilingual Evaluation of Generative AI”, in *Proceedings of the 2023 Conference on EMNLP*, Dec. 2023. DOI: 10.18653/v1/2023.emnlp-main.258
- [7] **R. Hada**, A. Ebrahimi Fard, S. Shugars, F. Bianchi, P. Rossini, D. Hovy, R. Tromble, and N. Tintarev, “Beyond Digital “Echo Chambers”: The Role of Viewpoint Diversity in Political Discussion”, in *Proceedings of the Sixteenth ACM International Conference on WSDM*, Feb. 2023. DOI: 10.1145/3539597.3570487
- [8] **R. Hada**, A. Seth, H. Diddee, and K. Bali, ““Fifty Shades of Bias”: Normative Ratings of Gender Bias in GPT Generated English Text”, in *Proceedings of the 2023 Conference on EMNLP*, Dec. 2023. DOI: 10.18653/v1/2023.emnlp-main.115
- [9] **R. Hada**, S. Sudhir, P. Mishra, H. Yannakoudakis, S. M. Mohammad, and E. Shutova, “Ruddit: Norms of Offensiveness for English Reddit Comments”, in *Proceedings of the 59th Annual Meeting of the ACL*, Aug. 2021. DOI: 10.18653/v1/2021.acl-long.210

PATENTS

- [10] N. Pareek, **R. Hada**, S. Malyala, and N. V. J. K. Kartik, “Synthetic Data Generation System and Method”, US Patent 12,608,610, 2026.
- [11] N. Pareek, **R. Hada**, and G. Sapra, “System for Detecting Errors in Generative AI Outputs”, US Patent 12,596,629, 2026.